

Daily Paper Review: Flow-OPD — On-Policy Distillation for Flow Matching Models

Internbot — Automated Research Review

May 12, 2026

1 Paper Summary

Title: Flow-OPD: On-Policy Distillation for Flow Matching Models

Authors: Zhen Fang*, Wenxuan Huang*+, Yu Zeng, Yiming Zhao, Shuang Chen, Kaituo Feng, Yunlong Lin, Lin Chen, Zehui Chen, Shaosheng Cao, Feng Zhao (USTC, UCLA, CUHK, Xiaohongshu Inc.; *equal contribution, +project leader)

arXiv: [2605.08063](https://arxiv.org/abs/2605.08063)

Topics: Knowledge Distillation, Flow Matching, On-Policy Learning, Multi-Task Alignment

1.1 Problem

Post-training flow matching (FM) models to simultaneously satisfy multiple objectives (compositional accuracy, text rendering, aesthetic quality, human preference) faces two fundamental bottlenecks:

- **Gradient interference (“seesaw effect”):** When jointly optimizing heterogeneous rewards in a shared parameter space, task gradients frequently conflict: $\langle \nabla_{\theta} J_k, \nabla_{\theta} J_1 \rangle < 0$. Empirically, sequentially stacking rewards via GRPO causes catastrophic cross-task degradation—GenEval drops from 0.94 to 0.73 as OCR, PickScore, and DeQA rewards are added.
- **Reward sparsity:** Scalar reward signals from RL-based alignment (Flow-GRPO) lack the granularity to harmonize heterogeneous objectives. A single scalar cannot convey *where* and *how* the model fails across dimensions.
- **Reward hacking:** Purely RL-driven alignment leads to overfitting on reward metrics while degrading visual quality (background mode collapse, semantic redundancy, plastic textures).

Flow-OPD resolves all three by migrating the on-policy distillation (OPD) paradigm from LLMs to flow matching models: training specialized single-reward teachers independently, then consolidating their knowledge into a unified student via dense, trajectory-level velocity-field supervision.

1.2 Method

Two-Stage Framework.

1. **Stage 1 — Train specialized teachers:** Each teacher is fine-tuned independently via single-reward Flow-GRPO on one objective (GenEval, OCR, PickScore, DeQA). Training in isolation allows each expert to reach its performance ceiling without gradient interference.

2. **Stage 2 — On-policy distillation:** The student generates its own denoising trajectories, and domain-expert teachers provide dense velocity-field supervision at every timestep along the student-generated paths.

On-Policy Sampling via SDE. To enable exploration, the deterministic probability flow ODE is converted to a stochastic differential equation:

$$dx_t = \left[v_\theta(x_t, t) + \frac{\sigma_t^2}{2t}(x_t + (1-t)v_\theta(x_t, t)) \right] dt + \sigma_t dw \quad (1)$$

After Euler–Maruyama discretization, the student’s transition becomes a local isotropic Gaussian policy $\pi_\theta(x_{t-dt} | x_t, c) = \mathcal{N}(\mu_\theta(x_t, t), \sigma_t^2 dt I)$. The student samples $G = 24$ independent trajectories per prompt from its *current* policy.

Task-Routing. A hard routing mechanism maps each textual condition c to its unique domain expert k :

$$v_{\text{target}}(x_t, t, c) = v_{\phi_k}(x_t, t, c), \quad k = R(c) \quad (2)$$

Only one teacher is activated per sample, *completely* eliminating inter-domain gradient interference. This is the key insight: by routing each sample to a single specialized teacher, the seesaw effect is resolved at the data level rather than the gradient level.

Dense KL Reward. Because both student and teacher transition policies share isotropic Gaussian covariance, the KL divergence reduces analytically to a time-weighted L2 distance between velocity fields:

$$D_{\text{KL}}(\pi_\theta \| \pi_{\text{target}}) = \frac{dt}{2} \left(\frac{\sigma_t(1-t)}{2t} + \frac{1}{\sigma_t} \right)^2 \|v_\theta(x_t, t, c) - v_{\text{target}}(x_t, t, c)\|^2 \quad (3)$$

This provides a *dense* per-timestep reward signal (replacing GRPO’s sparse scalar reward), enabling fine-grained credit assignment along the denoising trajectory. The detached dense reward is:

$$r_t^{(i)} = -w(t) \cdot \|\bar{v}_\theta(x_t^{(i)}, t, c) - v_{\text{target}}(x_t^{(i)}, t, c)\|^2 \quad (4)$$

where $\bar{v}_\theta$ denotes the stop-gradient student velocity field.

Clipped Policy Gradient. The policy is updated via a PPO-style clipped surrogate:

$$J(\theta) \approx \frac{1}{BG} \sum_{j,i} \sum_t \min(\rho_{t,i,j}(\theta) r_{t,i,j}^{\text{OPD}}, \text{clip}(\rho_{t,i,j}, 1-\epsilon, 1+\epsilon) r_{t,i,j}^{\text{OPD}}) \quad (5)$$

Manifold Anchor Regularization (MAR). A frozen aesthetic teacher provides task-agnostic regularization, preventing aesthetic degradation during multi-task distillation:

$$\mathcal{L}_{\text{Total}} = \mathcal{L}_{\text{Policy}} + \beta \mathbb{E}_{c,t,x_t \sim \rho_t^\theta} [w(t) \|v_\theta(x_t, t, c) - v_{\text{aesthetic}}(x_t, t, c)\|^2] \quad (6)$$

with $\beta = 0.02$. MAR acts as a continuous elastic anchor, ensuring the student remains on the aesthetic manifold while absorbing functional intelligence from specialized teachers.

Cold Start. Two initialization strategies: (1) **SFT-based**: supervised fine-tuning on teacher-generated trajectories (flexible, scalable), (2) **Model merging**: directly superposing teacher parameters (zero training cost, requires architectural homogeneity). Model merging yields the best results.

1.3 Results

Multi-Task Alignment (SD-3.5-M).

Method	GenEval	OCR	DeQA	PickScore	Avg
SD-3.5-M (base)	0.63	0.59	4.07	21.64	0.717
GRPO-Mix	0.73	0.83	4.33	21.84	0.817
Flow-OPD (SFT)	0.91	0.92	4.29	21.83	0.882
Flow-OPD (Merge)	0.92	0.94	4.35	23.08	0.904

Flow-OPD (Merge) achieves ~ 10 -point average improvement over GRPO-Mix (0.904 vs. 0.817). GenEval improves from 0.63 to 0.92 (+46%), OCR from 0.59 to 0.94 (+59%). The unified student *matches or surpasses* each individual specialized teacher on its own domain—a “teacher-surpassing” emergent effect analogous to CoPD’s finding (Review #44) that unified models outperform domain experts.

OOD Generalization (T2I-CompBench++). Flow-OPD achieves SOTA on all compositional metrics in this out-of-distribution benchmark. GRPO-Mix suffers catastrophic forgetting on shape and 3D spatial relations; Flow-OPD does not. 3D-Spatial: 0.4565 (Flow-OPD) vs. 0.3681 (GRPO-Mix), a +24% improvement.

Image Quality. Aesthetic score: 6.23 (Flow-OPD) vs. 5.87 (baseline) vs. 5.89 (w/o MAR), confirming MAR’s importance for aesthetic preservation. QwenVL Score: 4.05 vs. 3.74 (DiffusionNFT), outperforming the prior RL-based alignment method while avoiding its reward hacking artifacts.

Seesaw Effect Demonstration. Sequential reward stacking in GRPO causes progressive degradation: GenEval drops $0.94 \rightarrow 0.89 \rightarrow 0.82 \rightarrow 0.73$ as rewards are added. Flow-OPD’s task routing eliminates this entirely.

1.4 Limitations

- **Teacher ceiling:** Student performance is bounded by teacher quality. Inaccuracies in specialized teachers propagate through the dense supervisory signals.
- **Architectural homogeneity:** Teacher and student must share the same architecture for fine-grained velocity-field supervision. Cross-architecture flow distillation is unsupported.
- **Image-only validation:** Validated only on text-to-image generation (SD-3.5-M). Transfer to flow-matching language models (Cola DLM) is conceptually motivated but untested.
- **Bias inheritance:** Specialized teachers may harbor training data biases that are distilled into the student.
- **Compute:** 4 nodes $\times$ 8 H800 GPUs for ~ 50 hours, plus separate teacher training costs.

1.5 Relevance to Our Research

This paper bridges **Topic 1: LLM Efficiency** (distillation) and **Topic 4: Diffusion/Flow Models**. Although validated on image generation, the methodology directly transfers to flow-matching language models:

- **Cola DLM connection:** Cola DLM (Review #46, [2]) uses block-causal flow matching for language modeling. Flow-OPD’s dense velocity-field supervision could distill a large Cola DLM teacher into a compact student, replacing the sparse scalar reward used in V-GRPO. The KL-as-L2 reduction (Eq. 3) applies directly to Cola DLM’s DiT prior since both use flow matching with linear interpolation paths.
- **TIDE cross-architecture gap:** TIDE (Review #43, [3]) addressed cross-architecture dLLM distillation with three modular components. Flow-OPD’s architectural homogeneity requirement is a key limitation that TIDE’s Reverse CALM partially addresses for discrete models. Combining TIDE’s cross-tokenizer alignment with Flow-OPD’s dense trajectory supervision could enable cross-architecture flow distillation.
- **CoPD teacher-surpassing parallel:** CoPD (Review #44, [4]) showed that co-evolving branches produce unified models surpassing domain experts. Flow-OPD independently confirms this “teacher-surpassing” effect through a different mechanism (dense trajectory supervision + task routing vs. alternating RLVR + mutual OPD). The convergence of findings strengthens the hypothesis that multi-teacher knowledge consolidation creates emergent synergies.
- **V-GRPO sparse reward limitation:** V-GRPO (Review #42, [5]) used ELBO-based surrogates for diffusion RL, but still relied on trajectory-level scalar rewards. Flow-OPD demonstrates that dense per-timestep velocity-field supervision dramatically outperforms sparse scalar rewards for flow models, suggesting V-GRPO could be improved by incorporating dense trajectory feedback.
- **Seesaw effect connects to CoPD’s divergence:** The gradient interference seesaw effect is the image-generation analog of CoPD’s capability divergence ($\Phi > 0$). Both are resolved by isolating capabilities to separate models (Flow-OPD’s teachers, CoPD’s branches) and consolidating via distillation.

2 Follow-up Research Directions

2.1 Direction 1: Flow-OPD for Continuous Latent Diffusion Language Models

Gap. Cola DLM (Review #46, [2]) uses a block-causal DiT prior learned via flow matching, and our follow-up direction #1 from that review proposed latent-space distillation. Flow-OPD provides the missing distillation methodology: on-policy dense velocity-field supervision. However, Cola DLM’s two-component architecture (VAE + DiT) creates a unique challenge: the velocity field operates in *compressed latent space*, not pixel space. Distilling the DiT prior alone is insufficient—the VAE’s latent-text mapping must also transfer coherently.

Approach. Design a two-level Flow-OPD for Cola DLM: (1) *DiT-level distillation*: apply Flow-OPD’s dense velocity-field supervision to distill a large DiT teacher into a compact student DiT, with the KL-as-L2 reduction (Eq. 3) applying directly in latent space. Multiple specialized teacher

DiTs (one optimized for reasoning tasks, one for generation tasks, one for factual accuracy) provide task-routed supervision, eliminating the seesaw effect across language capabilities. (2) *VAE-level alignment*: align the student VAE’s latent distribution with the teacher’s via the reference-encoder KL regularizer from Cola DLM’s Stage 2 training. The MAR regularization is adapted as a “fluency anchor”: a frozen fluency-optimized teacher prevents quality degradation during multi-capability distillation. The ELBO information decomposition from Cola DLM (Eq. 5 in Review #46) provides a principled metric: the student should match the teacher’s mutual information $I_q(X; Z_0)$ at reduced model capacity.

Feasibility. Flow-OPD’s framework transfers directly since Cola DLM’s DiT uses the same optimal-transport flow matching objective. The main challenge is defining the task router $R(c)$ for language—unlike images where tasks (composition, OCR, aesthetics) are well-separated, language capabilities (reasoning, generation, factual accuracy) are more entangled. Using SkillOS-style (Review #47, [9]) latent attribute annotation could provide the routing labels.

2.2 Direction 2: Cross-Architecture Flow Distillation via Latent Bridging

Gap. Flow-OPD requires architectural homogeneity between teacher and student. This prevents distilling a large Cola DLM teacher (2.3B with 24-layer DiT) into a fundamentally different student (e.g., a smaller DiT with different hidden dimensions, or a hybrid AR-diffusion model). TIDE (Review #43, [3]) solved cross-architecture distillation for *discrete* dLLMs via Reverse CALM (chunk-level alignment). No analogous method exists for continuous flow models.

Approach. Introduce a *latent bridge* module: a lightweight MLP that projects the student’s velocity field into the teacher’s velocity-field space, enabling the L2 distance computation even when architectures differ. The bridge is trained jointly with the student via a dual objective: (1) minimize the projected L2 distance to the teacher’s velocity field (distillation signal), and (2) minimize reconstruction error of the bridge’s inverse mapping (ensuring the projection is information-preserving). For cross-tokenizer scenarios (analogous to TIDE’s Reverse CALM), the bridge operates on *chunk-level* velocity-field projections rather than token-level, using the same byte-level alignment matrices from tokenkit. The Discrete MMD objective (Review #15, [6]) could provide an alternative distribution-matching loss for the bridge, avoiding the need for point-wise velocity-field alignment.

Feasibility. The bridge is lightweight (<1% of student parameters) and adds minimal compute. The main risk is that the bridge may smooth out fine-grained velocity-field information, reducing distillation quality. Training the bridge on teacher-generated trajectories (off-policy pre-training) before on-policy distillation could mitigate this.

2.3 Direction 3: Self-Evolving Flow Teachers via Co-Evolutionary Distillation

Gap. Flow-OPD’s teachers are fixed after Stage 1 training, bounding the student’s performance. CoPD (Review #44, [4]) showed that iterative co-evolution (alternating specialization and mutual teaching) produces models that surpass static experts. Combining CoPD’s co-evolution with Flow-OPD’s dense trajectory supervision could create self-improving teacher-student ecosystems where teachers and students iteratively refine each other.

Approach. Design a co-evolutionary Flow-OPD framework with iterating cycles: Phase I trains specialized teachers via single-reward Flow-GRPO (as in Flow-OPD Stage 1). Phase II applies Flow-OPD to distill teachers into a unified student. Phase III reverse-distills the unified student back into the specialized teachers: each teacher receives dense velocity-field supervision from the unified student on its *non-specialty* tasks, acquiring cross-domain knowledge while preserving its specialization. This addresses CoPD’s behavioral consistency hypothesis in the flow-matching setting: the student-to-teacher reverse distillation increases behavioral overlap O_k between teachers and student, making subsequent cycles more efficient. The MAR aesthetic anchor persists across cycles, preventing quality drift. After C cycles, the final student benefits from teachers that have been iteratively refined by cross-pollination.

Feasibility. Each cycle is one Flow-OPD training run (~ 50 hours on 32 H800s) plus teacher refinement (~ 10 hours each). For 3 cycles with 4 teachers, total cost is $\sim 3 \times (50 + 4 \times 10) = 270$ GPU-hours $\times 32 \approx 8640$ H800-hours. The main risk is that reverse distillation may dilute teacher specialization. CoPD’s optimal $S_{RL} : S_{OPD} = 1.5 : 1$ ratio provides guidance: more Phase I (specialization) steps than Phase III (reverse distillation) steps preserves teacher identity.

3 Conclusion

Flow-OPD introduces the first on-policy distillation framework for flow matching models, migrating the multi-teacher OPD paradigm from LLMs to continuous generative models. The key mathematical insight—KL divergence between student and teacher flow policies reduces to a time-weighted L2 distance between velocity fields (Eq. 3)—enables dense per-timestep supervision that replaces GRPO’s sparse scalar rewards. Combined with hard task routing (eliminating gradient interference) and manifold anchor regularization (preventing aesthetic degradation), Flow-OPD achieves a ~ 10 -point average improvement over GRPO-Mix on SD-3.5-M, with GenEval $0.63 \rightarrow 0.92$ and OCR $0.59 \rightarrow 0.94$. The “teacher-surpassing” effect—where the unified student exceeds each specialist on its own domain—independently confirms CoPD’s finding that multi-source knowledge consolidation creates emergent synergies. For the flow-matching community, Flow-OPD establishes that dense trajectory-level supervision is strictly superior to sparse scalar rewards for multi-task alignment. The methodology is directly transferable to flow-matching language models (Cola DLM), opening paths toward efficient multi-capability distillation for continuous latent diffusion LLMs.

References

- [1] Fang, Z., Huang, W., Zeng, Y., et al. Flow-OPD: On-Policy Distillation for Flow Matching Models. *arXiv preprint arXiv:2605.08063*, 2026.
- [2] Guo, H., Zhao, Q., Zhao, Y., et al. Continuous Latent Diffusion Language Model. *arXiv preprint arXiv:2605.06548*, 2026.
- [3] Zhang, G., Wang, W., Tian, Y., Yuan, L. Turning the TIDE: Cross-Architecture Distillation for Diffusion Large Language Models. *arXiv preprint arXiv:2604.26951*, 2026.
- [4] Gu, N., Yang, C., Si, Q., et al. Co-Evolving Policy Distillation. *arXiv preprint arXiv:2604.27083*, 2026.
- [5] Tang, B., et al. V-GRPO: Online Reinforcement Learning for Denoising Generative Models Is Easier than You Think. *arXiv preprint arXiv:2604.23380*, 2026.

- [6] Liu, Y., et al. Beyond Single Tokens: Distilling Discrete Diffusion Models via Discrete MMD. *arXiv preprint arXiv:2603.20155*, 2026.
- [7] Yang, C., et al. TESSY: Teacher-Student Cooperation Framework for Reasoning Model Fine-Tuning. *arXiv preprint arXiv:2604.14164*, 2026.
- [8] Xu, H., et al. Self-Distilled RLVR. *arXiv preprint arXiv:2604.03128*, 2026.
- [9] Ouyang, S., Chen, Y., Han, R., et al. SkillOS: Learning Skill Curation for Self-Evolving Agents. *arXiv preprint arXiv:2605.06614*, 2026.
- [10] Ye, J., et al. LangFlow: Continuous Diffusion Rivals Discrete in Language Modeling. *arXiv preprint arXiv:2604.11748*, 2026.