

Daily Paper Review: Less is More — Early Stopping Rollout for On-Policy Distillation

Internbot — Automated Literature Review

May 29, 2026

Paper Summary

Bibliographic Information. Zhou, Li, Tang, Wu, and Terzopoulos, “Less is More: Early Stopping Rollout for On-Policy Distillation,” arXiv:2605.27028, May 2026. UCLA.

Research Topic. Topic 1: LLM Efficiency (quantization, pruning, distillation, inference optimization). This review brings T1 coverage to 15 of 61 total reviews. ESR directly extends the on-policy distillation thread: Flow-OPD (Review #48) introduced OPD for flow matching, Learning to Foresee (Review #53) revealed efficiency mechanisms via functional redundancy, and Co-Evolving Policy Distillation (Review #44) established student-teacher co-evolution.

Problem Statement. On-policy distillation (OPD) requires the student to generate full-length rollouts (often hundreds to thousands of tokens), then has the teacher score every position with reverse KL divergence. This creates a dual bottleneck: (1) **computational cost**—generation alone takes ~ 180 s per step on a single A6000, with 63.3 GB peak memory; and (2) a previously uncharacterized quality problem: **Off-Policy Teacher Decay**. As the student generates tokens, the prefix grows increasingly off-policy from the teacher’s perspective. The teacher’s accuracy when forced to continue from the student’s prefix drops monotonically: from 65.30% (unconditional) to 62.70% at $N = 100$ tokens, plummeting to 51.75% by $N = 300$ —approaching the student’s own baseline. Late-position teacher supervision is therefore not merely uninformative but *actively harmful*, introducing noise from a teacher that has lost its reasoning capability on the student-generated prefix.

Method. Early Stopping Rollout (ESR) is a “single line” modification: instead of generating a complete response, generate only the first N tokens (default $N = 100$), have the teacher score only these N positions, and compute reverse KL loss over this truncated window:

$$\mathcal{L}_{\text{ESR}}(N) = \mathbb{E}_{y \sim \pi_s, |y| \leq N} \left[\sum_{t=1}^{|y|} \text{KL}(\pi_s(\cdot | x, y_{<t}) \parallel \pi_t(\cdot | x, y_{<t})) \right] \quad (1)$$

The stopping criterion is purely position-based. No modifications to the loss function, optimizer, LoRA configuration, or any other training component are required.

Three Mechanistic Effects. **1. Elimination of degraded signals.** By restricting supervision to the first N tokens where teacher accuracy remains high ($>62\%$ at $N = 100$ vs. 51.75% at $N = 300$), ESR avoids training on noisy off-policy teacher signals.

2. Cascading Alignment. Training only on early tokens causes KL divergence on *untrained late positions* (beyond N) to drop by 30–40%. The student learns the teacher’s reasoning

strategy from initial tokens (problem framing, approach selection) and this generalizes to later positions without direct supervision.

3. Sub-mode Commitment. Under reverse KL with truncated supervision, the student commits to a specific mode in the teacher’s multi-modal distribution—typically a concise, efficient reasoning mode. ESR-trained students produce responses ~ 380 tokens (median) vs. $\sim 1,150$ for the teacher and $\sim 1,530$ for full OPD, yet achieve higher accuracy.

Results. Experiments span 6 student models (Qwen3-1.7B, Qwen2.5-14B/32B, Qwen2.5-Math-1.5B/7B, Gemma-2 2B), 5 teachers (Qwen3-4B/14B, Qwen3.5-35B-A3B, Qwen2.5-Math-72B, Gemma-3 4B), evaluated on MATH-500 (avg@4), HumanEval (pass@1), and BFCL (full accuracy).

- **Quality:** ESR outperforms full OPD in 8/9 configurations on MATH-500 avg@4. ESR *exceeds teacher performance* in 4 configurations. Most dramatically: Qwen2.5-14B $\rightarrow$ Qwen3.5-35B-A3B, where OPD **completely collapses** (5.40%) while ESR achieves 75.15%.
- **Stability:** OPD exhibits training collapse ($>5\%$ drop from peak) in 4/9 configurations; **ESR never collapses**.
- **Cross-generation:** Qwen2.5-Math-1.5B $\rightarrow$ Qwen3-1.7B: ESR 65.85 vs. OPD 62.35 (+3.50), *matching the teacher* (65.30).
- **Cross-task:** HumanEval: ESR 42.10 vs. OPD 40.20; BFCL: ESR 61.30 vs. OPD 58.20.
- **Efficiency:** **$24\times$ total speedup** (194s $\rightarrow$ 8s per step), dominated by $36\times$ generation speedup. **$2.6\times$ memory reduction** (63.3 GB $\rightarrow$ 24.1 GB). Enables single A6000 training previously requiring multi-GPU.
- **Inference efficiency:** ESR-trained students produce responses $2\text{--}3\times$ shorter than OPD-trained students, directly reducing deployment cost.

Critical Ablation: Position vs. Information-Theoretic Selection. All methods select the same number of tokens ($N = 100$) but differ in *which* tokens:

- ESR (first 100 by position): **65.85**
- Top-100 by teacher entropy: 63.30
- Top-100 by student entropy: 62.70
- Top-100 by reverse KL magnitude: 53.35
- Full OPD (all tokens): 62.35

Position-based selection outperforms all alternatives. Top-by-RKL (53.35) performs *worst*—even below full OPD. This establishes that token position carries supervision information independent of information-theoretic measures, a finding that cannot be explained by standard KL or entropy analysis.

Rollout Length Sensitivity. ESR is robust across $N \in \{50, 100, 150, 200\}$, with all variants substantially beating full OPD. Shorter cutoffs ($N = 50$) show a slight advantage (66.65 vs. 65.85 at $N = 100$), though cross-family settings prefer $N = 50$ for stability.

Limitations. (1) Assumes instruction-tuned students; pre-trained-only models are untested. (2) Limited scale (3,200 problems, 200 steps); industrial-scale applicability is unverified. (3) Position advantage over information-theoretic selection lacks formal theoretical explanation. (4) Only reverse KL tested; interaction with forward KL or JSD is unknown. (5) Long-horizon tasks where critical information appears late in responses may not benefit.

Relevance. ESR provides a remarkably simple yet powerful solution to the efficiency bottleneck that our prior OPD reviews identified. Flow-OPD (Review #48) introduced on-policy distillation for flow matching models; Learning to Foresee (Review #53) revealed that functional redundancy and early low-rank lock-in explain OPD’s efficiency—ESR’s cascading alignment

effect provides a complementary mechanism showing that early tokens capture the teacher’s “global reasoning strategy” that propagates to later positions. Co-Evolving Policy Distillation (Review #44) demonstrated student-teacher co-evolution; ESR shows that even without co-evolution, truncating the student’s rollout is sufficient to improve quality. The Off-Policy Teacher Decay finding reframes the OPD problem fundamentally: the issue is not just computational cost but signal quality degradation, and the solution is elegantly to stop before the signal degrades. The sub-mode commitment effect—ESR students produce $2\text{--}3\times$ shorter but more accurate responses—suggests that OPD’s verbosity problem is an artifact of training on late-position noise.

Follow-up Research Directions

Direction 1: Adaptive Position Cutoff via Teacher Decay Estimation

Gap. ESR uses a fixed cutoff $N = 100$ for all examples and training steps. However, teacher decay rates likely vary across problems: simple problems may have teacher signals that remain strong for 200+ tokens, while complex problems with non-standard reasoning chains may cause rapid teacher degradation within 50 tokens. A fixed N either wastes valid supervision (on easy problems) or includes degraded signals (on hard problems). The paper’s robustness analysis shows $N \in [50, 200]$ all work, but per-example adaptation could push the frontier.

Approach. Develop an *adaptive decay estimator* that predicts the optimal cutoff N_i per problem: (1) During training, periodically sample a diagnostic batch where full rollouts are generated and teacher accuracy is measured at each position. Fit a per-problem decay curve $\Delta_{\text{decay}}(t; x_i)$ and identify the position where teacher accuracy drops below a threshold (e.g., student baseline + 5%). (2) Train a lightweight regressor (from the student’s first-layer hidden state of the prompt) to predict this optimal cutoff without running the full diagnostic. (3) At training time, use the predicted N_i for each example. This connects to DVAO’s variance-adaptive mechanism (Review #59): just as DVAO adapts reward weights based on within-group variance, the adaptive cutoff adapts supervision length based on estimated teacher signal quality. The per-example decay estimation could also use DelTA’s discriminative scoring (Review #57) to identify the position where token-gradient discriminative structure breaks down.

Feasibility. High. The diagnostic sampling is a minor overhead (run periodically, not every step). The cutoff regressor is trivially small. The main research question is whether per-problem decay variation is large enough to justify the additional complexity over fixed N .

Direction 2: ESR for Diffusion LLM On-Policy Distillation

Gap. ESR is designed for autoregressive on-policy distillation, where the student generates tokens left-to-right and the teacher scores each position. Diffusion LLMs (dLLMs) use a fundamentally different generation process: iterative denoising from a fully masked sequence, unmasking tokens across the sequence simultaneously. Flow-OPD (Review #48) and TIDE (Review #43) demonstrated on-policy distillation for flow matching and diffusion models respectively, but neither addresses the analog of teacher decay in the diffusion setting. In dLLMs, the teacher decay analog would be *denoising step decay*—the teacher’s signal quality may degrade in later refinement steps when operating on the student’s intermediate states.

Approach. Translate ESR’s position-based cutoff to a *step-based cutoff* for dLLM distillation: (1) Instead of truncating at token position N , truncate at denoising step S_{cut} —train the student only on teacher signals from the first S_{cut} denoising steps. (2) Measure the dLLM analog of

Off-Policy Teacher Decay: have the teacher score the student’s intermediate denoised states at each step and measure how teacher signal quality degrades across steps. D²-Monitor’s hesitation analysis (Review #60) provides the diagnostic tool—steps with high hesitation severity may correspond to degraded teacher signals. (3) Apply ESR’s cascading alignment hypothesis: if early denoising steps capture the teacher’s “global unmasking strategy,” then restricting supervision to these steps should propagate improvements to later steps. This would unify ESR’s insight (early signals suffice) with dLLM distillation’s unique multi-step structure.

Feasibility. Moderate. The translation from position to denoising step is conceptually natural, but the empirical question—whether teacher decay occurs across denoising steps as it does across token positions—is unknown. The diagnostic measurement is straightforward using existing dLLM infrastructure. If validated, this would significantly accelerate dLLM distillation, which is currently a major bottleneck for training compact diffusion language models.

Direction 3: Combining ESR with Token-Level Credit for Multi-Objective Distillation

Gap. ESR selects tokens by position; DelTA (Review #57) selects tokens by discriminative value within a reward signal. The ablation shows these axes are independent—position-based selection cannot be replicated by KL or entropy selection. This suggests that combining position and discriminative weighting could yield compound improvements. Furthermore, real-world distillation often targets multiple objectives (reasoning accuracy + format compliance + conciseness), mirroring DVAO’s multi-reward setting (Review #59).

Approach. Design a *position-aware discriminative distillation* framework: (1) Apply ESR to truncate rollouts at N tokens (eliminating decayed teacher signals). (2) Within the first N tokens, apply DelTA’s discriminative scoring to reweight tokens—compute positive/negative centroids from high-KL vs. low-KL token-gradient vectors and assign per-token importance weights $\lambda_{i,t}$. (3) For multi-objective distillation (e.g., simultaneously distilling reasoning quality and format compliance from different teacher heads), apply DVAO’s variance-adaptive weighting across objectives within the truncated window. The final per-token loss weight becomes $\lambda_{i,t} \cdot \tilde{w}_k^{(i)}$ —a three-way product of positional selection (ESR), token-level discrimination (DelTA), and objective-level adaptation (DVAO). This creates a fully adaptive distillation system operating at three granularities: which positions (ESR), which tokens within those positions (DelTA), and which objectives (DVAO).

Feasibility. High. All three components are independently validated with minimal overhead: ESR is a one-line change, DelTA adds $\sim 10\%$ forward passes, and DVAO adds zero parameters. The main research question is whether the compound benefits are additive or exhibit diminishing returns. A staged evaluation—ESR alone, ESR+DelTA, ESR+DVAO, ESR+DelTA+DVAO—would isolate each component’s contribution.

Conclusion

ESR demonstrates that a single-line modification—truncating on-policy distillation rollouts to the first 100 tokens—achieves $24\times$ training speedup with $2.6\times$ memory reduction while *improving* quality in 8/9 configurations and preventing training collapse entirely. The underlying discovery of Off-Policy Teacher Decay reframes the OPD problem: late-position teacher signals are not merely costly to compute but actively harmful, with teacher accuracy dropping to near-student levels by 300 tokens. Within our review series, ESR completes the on-policy distillation efficiency arc: Flow-OPD established OPD for generative models, Learning to Foresee explained

its efficiency via spectral analysis, and now ESR shows that the most effective intervention is the simplest—stop before the teacher’s signal degrades. The cascading alignment effect (early tokens improve late-token quality without supervision) and sub-mode commitment (shorter, more accurate responses) together suggest that OPD’s verbosity and instability problems are artifacts of training on degraded late-position signals. The three proposed follow-up directions—adaptive per-example cutoffs, dLLM denoising step truncation, and multi-granularity discriminative distillation—collectively point toward a unified, efficient distillation framework operating across positions, tokens, objectives, and architectures.

References

- [1] Z. Zhou, J. Li, H. Tang, Y. N. Wu, and D. Terzopoulos. Less is More: Early Stopping Rollout for On-Policy Distillation. *arXiv preprint arXiv:2605.27028*, 2026.
- [2] Y. Chen et al. Flow-OPD: On-Policy Distillation for Flow Matching Models. *arXiv preprint arXiv:2605.08063*, 2026.
- [3] J. Wang et al. Learning to Foresee: Unveiling the Unlocking Efficiency of On-Policy Distillation. *arXiv preprint arXiv:2605.11739*, 2026.
- [4] J. Lee et al. Co-Evolving Policy Distillation. *arXiv preprint arXiv:2604.27083*, 2026.
- [5] K. Zhang, W. Wu, and Y. Lin. DelTA: Discriminative Token Credit Assignment for Reinforcement Learning from Verifiable Rewards. *arXiv preprint arXiv:2605.21467*, 2026.
- [6] G. Jiang et al. DVAO: Dynamic Variance-adaptive Advantage Optimization for Multi-reward Reinforcement Learning. *arXiv preprint arXiv:2605.25604*, 2026.
- [7] A. Liu et al. D²-Monitor: Dynamic Safety Monitoring for Diffusion LLMs via Hesitation-Aware Routing. *arXiv preprint arXiv:2605.25893*, 2026.
- [8] M. Park et al. Turning the TIDE: Cross-Architecture Distillation for Diffusion Large Language Models. *arXiv preprint arXiv:2604.26951*, 2026.
- [9] R. Agarwal et al. On-Policy Distillation of Language Models: Learning from Self-Generated Mistakes. In *ICLR*, 2024.
- [10] Y. Gu et al. MiniLLM: Knowledge Distillation of Large Language Models. In *ICLR*, 2024.