

Daily Paper Review: APPO — Agentic Procedural Policy Optimization

Internbot — Automated Research Summary

June 16, 2026

1 Paper Summary

1.1 Problem and Motivation

Agentic RL based on RLVR provides feedback only at the *trajectory level*—a single outcome reward for an entire multi-step interaction—creating a fundamental credit assignment problem. Wang et al. [1] argue that existing approaches attempt to extract finer credit by branching trajectories at intermediate locations, but rely on *coarse-grained heuristic units*: workflow-level approaches (ReAct, Reflexion) treat entire stages as atomic, while tool-call-level approaches (ARPO, Tree-GRPO) branch only at tool-call boundaries, compressing all reasoning between tool calls into monolithic “thinking” blocks. A pilot study reveals that (1) the highest-uncertainty positions are *not* concentrated at tool-call boundaries but are broadly distributed throughout the thinking span; (2) token entropy alone does not reliably indicate decision significance (high-entropy tokens may reflect lexical rarity rather than task-relevant choices); (3) the paper’s proposed Branching Score (BS) correlates much better with actual outcome divergence than entropy-based selection.

1.2 Method

APPO Architecture: Three Phases. *Phase 1—Initialization:* Given input x and global rollout budget M , generate N full rollouts from behavior policy π_{old} as tree roots.

Phase 2—Procedural Rollout Branching: For each token $H_{n,i}$ in rollout H_n , compute two quantities:

Future Value (Ω): The accumulated decayed importance sampling ratio measuring downstream distributional shift:

$$\Omega_{n,i} = \exp \left(\sum_{i' \geq i} \gamma^{i'-i} \log \frac{\pi_{\theta}(H_{n,i'} \mid H_{n,<i'}, x)}{\pi_{\text{old}}(H_{n,i'} \mid H_{n,<i'}, x)} \right) \quad (1)$$

where $\gamma = 2^{-1/\tau}$ with $\tau=32$ suppresses distant tokens. Large Ω indicates the policy has shifted its downstream behavior at that reasoning juncture.

Branching Score (BS): The *multiplicative* combination of z-score-normalized, clipped Ω and z-score-normalized entropy:

$$\text{BS}_{n,i} = Z(\text{clip}(\Omega_{n,i}, 1-\epsilon', 1+\epsilon'); H_n) \cdot Z(\mathcal{H}_{n,i}; H_n) \quad (2)$$

The multiplicative form ensures selected tokens are *both locally uncertain AND consequential downstream*—entropy alone selects rare nouns (vocabulary long-tail), while Ω alone selects confirmation tokens with high influence but no training value. Top- B tokens per rollout are selected, and branches are generated from the *current policy* π_{θ} (not π_{old}).

Phase 3—Procedural Advantage Estimation and Optimization:

Dual-Group Advantage: Since initial rollouts come from π_{old} and branches from π_{θ} , mixing them introduces distributional bias. APPO normalizes rewards *separately* within each group, then averages:

$$A_{n,i}^{\text{base}} = \text{avg} \left\{ \frac{R_n - \bar{R}_{T_*}}{\sigma_{T_*}} \mid T_* \in \{T_{\text{init}}, T_{\text{branch}}\} \right\} \quad (3)$$

Future-Aware Advantage Scaling: Amplifies credit at high-impact decision points:

$$A_{n,i} = A_{n,i}^{\text{base}} \cdot (1 + b \cdot A_{n,i}^{\text{fut}}) \quad (4)$$

where $A_{n,i}^{\text{fut}} = \text{clip}_{\epsilon'}(\Omega_{n,i})$ and $b=0.5$ controls the contribution. The final objective is a PPO-style clipped surrogate with KL regularization.

Theoretical Guarantees. *Theorem 1 (Variance Reduction):* Under a monotonicity assumption between BS and conditional reward variance, APPO’s BS-guided allocation achieves $\text{Var}(g_{\text{APPO}}) \leq \text{Var}(g_{\text{base}}) - \Delta_{\Omega}(\text{BS})$ where $\Delta_{\Omega} \geq 0$, via Lagrange-optimal sample allocation.

Theorem 2 (Policy Improvement Bound): Under KL and clipping constraints, $J(\pi_{\text{new}}) - J(\pi_{\text{old}}) \geq L_{\text{APPO}}(\pi_{\text{new}}) - C\epsilon/(1-\gamma)^2$, where L_{APPO} is the $\omega(s)$ -weighted surrogate that focuses improvement on procedurally important states.

1.3 Results

Deep Reasoning (10 benchmarks). On Qwen2.5-7B: APPO achieves **62.2%** average (+8.9% over GRPO’s 56.5%, +3.9% over ARPO’s 58.3%). On Llama3.1-8B: **57.4%** average (+7.9% over GRPO). Key gains vs. ARPO on Qwen2.5-7B: AIME24 +6.7pp, HotpotQA +7.7pp, 2Wiki +5.4pp, Bamboogle +6.1pp.

Deep Search (GAIA, HLE, Xbench). On Qwen3-8B: APPO achieves **42.7** GAIA average, **33.8** HLE, **13.4** Xbench (+3.9, +1.8, +4.6 over ARPO). On Qwen3-14B: **46.6** GAIA, **43.4** HLE, **12.3** Xbench. Notably, APPO with Qwen3-14B (46.6 GAIA) massively exceeds DeepSeek-R1-671B (25.2 GAIA)—a model $\sim 48\times$ larger.

Budget Efficiency. At budget $M = 4$, APPO matches ARPO at $M = 16$ (52.7 vs. 52.7 on knowledge-intensive tasks), demonstrating **4 $\times$ budget efficiency**. Balanced allocation ($N=4, B=3$ at $M=16$) outperforms both extreme breadth ($N=8, B=1$) and extreme depth ($N=2, B=7$).

Ablation. All three components contribute non-redundant gains. Average degradation when removed (across both backbones): A^{fut} most impactful (−1.4 to −3.4pp), dual-group advantage (−1.1 to −2.1pp), BS vs. entropy (−0.9 to −1.8pp).

1.4 Limitations

No theoretical optimality guarantee for BS as a branching criterion. Limited tool diversity (search and Python only). Scale behavior at 32B+ untested. KL regularization disabled ($\beta=0$) throughout training, raising questions about the applicability of the policy improvement bound. Computational overhead of BS computation (requiring forward passes through both π_{θ} and π_{old}) not quantified. Dependence on ARPO’s strong SFT initialization pipeline.

1.5 Relevance to Our Research

APPO introduces *procedure-level* credit assignment for agentic RL, occupying a semantic middle ground between the token-level analysis we reviewed in DelTA [2] and the trajectory-level GRPO approaches in BandPO [3] and DVAO [4]. The Branching Score—which multiplies future value (Ω , measuring downstream distributional shift) by local entropy—operationalizes the “sparse and critical” token insight from our geometric OPD review [5], where we saw that OPD’s update subspace locks early to a low-dimensional channel. APPO’s Ω formulation essentially identifies tokens that *break* this lock—positions where the policy’s downstream trajectory diverges most from its previous behavior.

The dual-group advantage design connects to Flow-DPPO’s [6] asymmetric masking: both recognize that mixing on-policy and off-policy data requires careful distributional separation. APPO’s $4\times$ budget efficiency through principled branching parallels the finding from ESR [7] and FiRe-OPD [8] that targeted sample selection dramatically reduces computational waste in on-policy methods. The future-aware advantage scaling (A^{fut}) provides a concrete mechanism for the credit assignment problem identified in our cross-domain RL interference review [9]: amplifying gradients at decision points where interference is most consequential.

2 Follow-up Research Directions

2.1 Direction 1: Procedural Branching for On-Policy Distillation

Gap. Current OPD methods (FiRe-OPD [8], ESR [7]) generate student rollouts and apply teacher supervision uniformly or with token-level reweighting. Our geometric OPD review [5] showed token sparsification preserves the locked update subspace, but the selection criteria (top-KL, random) are heuristic. APPO’s BS provides a principled alternative: identify tokens where alternative continuations would diverge most, and concentrate teacher supervision there.

Approach. Adapt APPO’s Branching Score to OPD: compute Ω as the accumulated student-teacher KL ratio (rather than $\pi_\theta/\pi_{\text{old}}$), multiply by local entropy, and allocate teacher forward passes only at high-BS positions. At these positions, generate branch continuations from the student and compute teacher KL supervision. For low-BS positions, use a cheaper proxy (cached teacher logits or no supervision). This creates a *procedure-aware distillation* scheme where teacher compute scales with the decision complexity of each token, potentially reducing teacher inference costs by 60–80% while matching full-supervision quality.

Feasibility. The Ω computation reuses the student’s log-probabilities already available during rollout. BS computation adds one z-score normalization per rollout. Teacher forward passes at high-BS positions only require caching the student rollout prefix. Implementation on Qwen3-8B (student) / Qwen3-32B (teacher) with DAPO-Math-17k, directly comparable to the geometric OPD experimental setup. Estimated timeline: 2–3 weeks.

2.2 Direction 2: Branching Score for Diffusion LLM Denoising Credit Assignment

Gap. RL fine-tuning of diffusion LLMs (V-GRPO [10], DARE [11]) applies trajectory-level rewards to iterative denoising processes, facing the same credit assignment problem APPO addresses for autoregressive agents. Each denoising step is a “decision point” where the model unmask tokens, but current methods assign equal credit to all steps. Flow-DPPO [6] introduced per-step divergence masking but did not identify which denoising steps are most consequential for the final output quality.

Approach. Adapt BS to the diffusion denoising MDP: at each reverse step t , compute Ω_t as the accumulated importance ratio across future denoising steps, measuring how much the current policy’s remaining trajectory differs from the reference. Multiply by per-step entropy (which is naturally high at early denoising steps where many tokens are masked). The resulting BS_t identifies “critical denoising steps” where branching would produce the largest outcome divergence. Use this for: (a) targeted GRPO branching in V-GRPO training (branch at high-BS steps rather than uniformly), (b) future-aware advantage scaling (A_t^{fut}) to amplify credit at these steps. The hypothesis is that early denoising steps (high uncertainty, many masked tokens) will have high BS, while later refinement steps will have low BS—providing natural adaptive credit allocation.

Feasibility. The denoising MDP has $K = 10\text{--}50$ steps (vs. hundreds of tokens in autoregressive rollouts), making Ω computation trivial. BS requires two forward passes per step (current and old policy), which are already performed in V-GRPO. Branch generation at high-BS denoising steps reuses the existing reverse sampling infrastructure. Validation on LLaDA-8B with AIME-24/MATH-500 using V-GRPO + BS-guided branching. Estimated compute: comparable to standard V-GRPO training.

2.3 Direction 3: Self-Evolving Agents with Procedure-Level Skill Memory

Gap. Our self-evolving agent reviews (SkillOS [12], SkillOpt [13], Continual Experience Internalization [14]) and EvoArena [15] identify two complementary challenges: (a) skill accumulation (learning reusable procedures from experience) and (b) memory evolution (tracking what changed and what persists). APPO’s procedure concept—latent decision points identified by BS—provides a natural unit for skill extraction. Rather than accumulating entire trajectories or workflow stages as skills, extract the *high-BS segments* as reusable procedural patterns that represent the agent’s learned decision-making at critical junctures.

Approach. After APPO training, for each successful trajectory, extract segments around the top- k BS positions (prefix context + decision token + continuation). Cluster these segments by semantic similarity to form a *procedural skill library*. Each skill entry includes: the decision context (what situation triggers this procedure), the chosen action (what the agent decided), the outcome impact (Ω magnitude), and EvoMem-style version patches when skills are updated. At inference, retrieve relevant procedural skills as in-context exemplars when the agent encounters similar decision contexts. This creates a bridge between APPO’s training-time procedure identification and deployment-time skill reuse, with EvoArena-style version tracking for robustness under environment evolution.

Feasibility. BS computation is already available from APPO training—extracting high-BS segments requires only post-processing of training logs. Semantic clustering can use embedding similarity (e.g., UMAP + DBSCAN as in APPO’s own diversity analysis). The procedural skill library is lightweight (one segment per decision pattern, not full trajectories). Integration with existing agent frameworks (OpenHands, Terminus2) requires adding a retrieval step at high-entropy inference positions. Validation on EvoArena’s Terminal-Bench-Evo (chain-level accuracy as metric) would directly test whether procedural skills improve consistency across evolving environments.

3 Conclusion

APPO introduces procedure-level credit assignment for agentic RL by identifying latent decision points through the Branching Score—a multiplicative combination of local entropy and

future value (Ω) that selects tokens both locally uncertain and consequential downstream. The three-phase architecture (initialization, BS-guided branching, dual-group advantage with future-aware scaling) achieves state-of-the-art results across 13 benchmarks spanning math reasoning, knowledge-intensive QA, and deep search, with $4\times$ budget efficiency over ARPO. The theoretical framework provides variance reduction and policy improvement guarantees grounding the BS-guided branching in trust-region optimization. For our research program, APPO establishes “procedures” as the natural granularity for agentic RL—finer than tool-call stages but coarser than individual tokens—and the BS scoring mechanism provides a principled bridge between our token-level credit assignment work (DelTA) and trajectory-level optimization (GRPO variants), with direct applications to on-policy distillation, diffusion LLM denoising, and self-evolving agent skill extraction.

References

- [1] Wang, X., Ma, Z., Wang, Y., Ji, Y., Yang, S., Chen, G., Wang, P., and Chu, X. “APPO: Agentic Procedural Policy Optimization.” *arXiv preprint arXiv:2606.12384*, 2026.
- [2] Wang, Z., et al. “DelTA: Discriminative Token Credit Assignment for Reinforcement Learning from Verifiable Rewards.” *arXiv preprint arXiv:2605.21467*, 2026.
- [3] Li, Y., et al. “BandPO: Bridging Trust Regions and Ratio Clipping via Probability-Aware Bounds for LLM Reinforcement Learning.” *arXiv preprint arXiv:2603.04918*, 2026.
- [4] Chen, Y., et al. “DVAO: Dynamic Variance-adaptive Advantage Optimization for Multi-reward Reinforcement Learning.” *arXiv preprint arXiv:2605.25604*, 2026.
- [5] Shen, Z., et al. “On the Geometry of On-Policy Distillation.” *arXiv preprint arXiv:2606.07082*, 2026.
- [6] Ping, B., et al. “Flow-DPPO: Divergence Proximal Policy Optimization for Flow Matching Models.” *arXiv preprint arXiv:2606.11025*, 2026.
- [7] Wang, Z., et al. “Less is More: Early Stopping Rollout for On-Policy Distillation.” *arXiv preprint arXiv:2605.27028*, 2026.
- [8] Li, Y., et al. “Filter, Then Reweight: Rethinking Optimization Granularity in On-Policy Distillation.” *arXiv preprint arXiv:2606.02684*, 2026.
- [9] Liu, J., et al. “A Local Perturbation Theory for Cross-Domain Interference and Recovery in Multi-Domain RL.” *arXiv preprint arXiv:2606.02398*, 2026.
- [10] Li, Y., et al. “V-GRPO: Online Reinforcement Learning for Denoising Generative Models Is Easier than You Think.” *arXiv preprint arXiv:2604.23380*, 2026.
- [11] Xu, H., et al. “DARE: Diffusion Large Language Models Alignment and Reinforcement Executor.” *arXiv preprint arXiv:2604.04215*, 2026.
- [12] Chen, X., et al. “SkillOS: Learning Skill Curation for Self-Evolving Agents.” *arXiv preprint arXiv:2605.06614*, 2026.
- [13] Li, Y., et al. “SkillOpt: Executive Strategy for Self-Evolving Agent Skills.” *arXiv preprint arXiv:2605.23904*, 2026.
- [14] Liu, J., et al. “Rethinking Continual Experience Internalization for Self-Evolving LLM Agents.” *arXiv preprint arXiv:2606.04703*, 2026.
- [15] Xu, J., et al. “EvoArena: Tracking Memory Evolution for Robust LLM Agents in Dynamic Environments.” *arXiv preprint arXiv:2606.13681*, 2026.