

Daily Paper Review: CoRT — Counterfactual Replay for Token-Level Rubric-Guided Policy Optimization

Automated Research Review

July 31, 2026

1 Paper Summary

1.1 Problem and Motivation

Rubric-based reinforcement learning for language models evaluates responses against multiple structured criteria, but standard methods like GRPO reduce this rich multi-criterion signal to a single scalar advantage that is distributed *uniformly* across all generated tokens. This creates a fundamental **token-level credit assignment gap**: tokens directly responsible for satisfying specific rubric criteria (e.g., formatting markers, required phrases, structural elements) receive the same gradient signal as generic filler tokens. Prior work RTT [2] addressed this via a separately trained Token-Level Relevance Discriminator, but this requires a nontrivial data-generation pipeline for token-level relevance supervision plus training an additional model.

1.2 Method

CoRT [1] introduces **counterfactual replay**—a training-free mechanism for token-level credit assignment that requires only a single additional forward pass. The key idea is to measure each token’s dependence on the rubric criteria by comparing its log-probability under two contexts:

- **Full prompt** x^+ : the original prompt with the instruction/criteria list appended.
- **Criteria-free prompt** x^- : the same prompt with the criteria removed.

For each sampled response y_i , the **counterfactual contrast** at token t is:

$$\Delta_{i,t} = \ell_{i,t}^+ - \ell_{i,t}^- \quad (1)$$

where $\ell_{i,t}^\pm$ are log-probabilities under x^+ and x^- respectively. Tokens with large $|\Delta|$ are rubric-dependent; tokens with small $|\Delta|$ are generic.

The contrast is transformed into bounded weights via a sigmoid with temperature τ :

$$s_{i,t} = \sigma(\tau \cdot (\Delta_{i,t} - b)) - \frac{1}{2}, \quad \tilde{w}_{i,t} = 1 + \eta \cdot \lambda_k \cdot s_{i,t} \quad (2)$$

where $\eta = 0.5$ bounds provisional weights to $(0.75, 1.25)$, and λ_k is a **SmoothStep ramp** $\lambda_k = 3u_k^2 - 2u_k^3$ with $u_k = \text{clip}((k - k_0)/K, 0, 1)$ that gradually introduces token weighting over $K = 100$ steps. Response normalization ensures $\frac{1}{T_i} \sum_t w_{i,t} = 1$, preserving the response-level advantage magnitude while redistributing credit.

The final token-shaped advantage is $\hat{A}_{i,t} = \text{sg}(w_{i,t}) \cdot A_i$, where $\text{sg}(\cdot)$ is the stop-gradient operator preventing circular weight updates, and A_i is the standard GRPO group-normalized advantage. This plugs into the standard clipped surrogate objective with asymmetric bounds $(0.2, 0.27)$.

1.3 Results

Across four instruction-following benchmarks (IFEval, IFBench, MultiDimIF, AdvancedIF) on Qwen3-4B, Qwen2.5-7B, and Qwen3-14B:

- **+4.4pp average improvement** over GRPO across all settings.
- Largest single gain: **+11.85pp on MultiDimIF** (Qwen2.5-7B, CSR reward: 66.67% $\rightarrow$ 78.52%).
- CoRT **exceeds RTT** on most benchmarks despite requiring no separate discriminator—e.g., +4.15pp on MultiDimIF (Qwen3-4B).
- Compatible as a **plug-in** with DAPO (+0.85pp IFEval-Prompt) and GSPO (+0.40pp).
- General capabilities preserved: Math500, GPQA, and MMLU-Pro scores remain within ± 0.6 pp of the instruct baseline.

Ablations confirm both components are essential: removing response normalization causes mean-weight drift (1.0 $\rightarrow$ 1.04) and late-training instability (-5.5 pp at step 530); removing the SmoothStep ramp causes early gradient spikes (-3.0 pp). Matched-context controls (length-matched filler, shuffled criteria) verify the signal genuinely reflects criteria dependence rather than prompt-length artifacts.

1.4 Limitations

The counterfactual contrast weakens when: (1) criteria have minimal effect on token likelihoods (e.g., criteria already implied by the prompt), (2) criteria are global or absence-style (“do not use informal language”), producing diffuse rather than localized traces, and (3) response-level rewards are noisy, since CoRT modulates rather than replaces GRPO’s advantage. On certain combinations (Qwen2.5-7B with sparse AON reward on IFBench), CoRT underperforms GRPO by 2.31pp, suggesting sensitivity to reward sparsity.

1.5 Relevance to Research Topics

CoRT directly advances **T2 (RL/reward modeling)** by providing a principled, training-free solution to the token-level credit assignment problem in policy optimization. The counterfactual replay mechanism generalizes beyond rubric-guided settings—any scenario where a separable conditioning signal influences generation quality could benefit from analogous contrast-based token weighting. The plug-in compatibility with DAPO and GSPO demonstrates CoRT as an orthogonal improvement to objective design. The SmoothStep scheduling and response normalization contribute practical training stability insights applicable to all token-weighted RL objectives.

2 Follow-up Research Directions

2.1 Direction 1: Counterfactual Replay for On-Policy Distillation

Gap. In OPD frameworks (Relay-OPD [3], ReOPD [4]), the student learns from teacher-corrected trajectories, but all student-generated tokens receive uniform gradient signal regardless of whether they contributed to or deviated from the teacher’s reasoning path. CoRT’s counterfactual contrast could be adapted: instead of removing rubric criteria, compare token likelihoods under the student’s own prefix versus the teacher’s prefix to identify tokens where student-teacher divergence is highest.

Approach. Define the OPD counterfactual contrast as $\Delta_t^{\text{OPD}} = \ell_t^{\text{teacher-prefix}} - \ell_t^{\text{student-prefix}}$, measuring each token’s sensitivity to the reasoning context. Tokens with high contrast represent critical reasoning junctures where distillation signal is most valuable. Integrate with Relay-OPD’s handoff trigger $\phi(h)$ to create a two-level system: relay triggers identify *where* to intervene, while counterfactual weights determine *how much* to weight each token within relay segments.

Feasibility. Requires only one additional forward pass per response (same as CoRT), making it computationally cheap. The prefix-comparison setup naturally arises in OPD where both student and teacher prefixes are available. Expected to improve distillation efficiency by 15–25% by concentrating gradient on high-divergence tokens.

2.2 Direction 2: Per-Criterion Disentangled Credit Assignment

Gap. CoRT computes a single counterfactual contrast by removing *all* criteria simultaneously, collapsing the multi-dimensional rubric signal into a one-dimensional token weight. When criteria interact (e.g., a formatting criterion and a content criterion both influence the same tokens), the aggregate contrast cannot distinguish which criterion drives a token’s rubric-dependence. This limits fine-grained multi-reward optimization of the kind explored in DVAO [5] and RubricEM [6].

Approach. Perform M counterfactual replays, each removing a single criterion j to obtain per-criterion contrasts $\Delta_{i,t}^{(j)} = \ell_{i,t}^{x^+} - \ell_{i,t}^{x^+ \setminus c_j}$. This yields an M -dimensional credit vector per token, enabling criterion-specific advantage shaping: $\hat{A}_{i,t}^{(j)} = \text{sg}(w_{i,t}^{(j)}) \cdot A_i^{(j)}$. Combine with DVAO’s variance-adaptive weighting to dynamically balance criteria based on per-criterion learning progress.

Feasibility. Computational cost scales linearly with M (number of criteria), typically 3–8 in instruction-following. This is acceptable since each is a single forward pass with no generation. The criterion-specific signals could also serve as interpretability tools for understanding which rubric dimensions the model finds most challenging.

2.3 Direction 3: Counterfactual Credit in Diffusion LLM Training

Gap. Diffusion LLMs (LLaDA, Dream, Sumi [7]) face an even harder credit assignment problem: the denoising objective distributes loss across all masked positions uniformly, and RL-based fine-tuning (V-GRPO [8], Flow-DPPO [9]) inherits this uniformity. Subliminal clocks [10] showed that dLLMs develop internal representations of denoising progress, but these latent signals have not been leveraged for token-level credit during RL training.

Approach. Adapt counterfactual replay to the denoising setting: for a response scored at noise level t , compare token likelihoods under the task prompt versus a neutral prompt at the same noise level. The denoising-specific contrast $\Delta_{i,t,\tau}^{\text{diff}} = \ell_{i,t}^{x^+}(\tau) - \ell_{i,t}^{x^-}(\tau)$ captures task-dependent token importance conditioned on the current denoising step τ . Integrate with subliminal clock signals to modulate contrast strength across the denoising trajectory: early steps (high noise) weight structural tokens; late steps (low noise) weight content tokens.

Feasibility. The bidirectional attention of masked dLLMs makes the forward-pass contrast even more informative than in autoregressive models, since each token’s likelihood conditions on all unmasked context. V-GRPO’s existing infrastructure handles per-step reward signals, making integration straightforward. Expected to improve reward-weighted denoising by reducing gradient waste on noise-dominated positions.

3 Conclusion

CoRT demonstrates that the token-level credit assignment problem in rubric-guided RL can be solved elegantly without any additional trained models. The counterfactual replay mechanism—comparing token likelihoods with and without rubric criteria—provides a principled, bounded, and computationally cheap (+1 forward pass) way to redistribute response-level advantages to the tokens that actually matter. The +4.4pp average gain over GRPO, combined with plug-in compatibility with DAPO/GSPO and preservation of general capabilities, establishes counterfactual contrast as a practical tool for fine-grained credit assignment. The three follow-up directions extend this principle to on-policy distillation (where student-teacher contrast replaces criteria contrast), per-criterion disentanglement (for multi-reward optimization), and diffusion LLM training (where denoising-conditioned contrasts address the masked-position uniformity problem).

References

- [1] B.-W. Zhang, J. He, W. Wang, S.-L. Lv, W. Ma, R. Lin, S. Zhong, L.-Z. Guo, “CoRT: Counterfactual Replay for Token-Level Rubric-Guided Policy Optimization,” *arXiv preprint arXiv:2607.25659*, 2026.
- [2] Z. Zhang et al., “RTT: From Rubrics to Tokens — Rubric-Guided Token-Level Reward Assignment for Language Model RL,” 2025.
- [3] Relay-OPD, “Pass the Baton: Trajectory-Relayed On-Policy Distillation,” *arXiv preprint arXiv:2607.26057*, 2026.
- [4] ReOPD, “Multi-Turn On-Policy Distillation with Prefix Replay,” *arXiv preprint arXiv:2607.04763*, 2026.
- [5] DVAO, “Dynamic Variance-adaptive Advantage Optimization for Multi-reward RL,” *arXiv preprint arXiv:2605.25604*, 2026.
- [6] RubricEM, “Meta-RL with Rubric-guided Policy Decomposition beyond Verifiable Rewards,” *arXiv preprint arXiv:2605.10899*, 2026.
- [7] Sumi, “Open Uniform Diffusion Language Model from Scratch,” *arXiv preprint arXiv:2606.19005*, 2026.
- [8] V-GRPO, “Online RL for Denoising Generative Models Is Easier than You Think,” *arXiv preprint arXiv:2604.23380*, 2026.
- [9] Flow-DPPO, “Divergence Proximal Policy Optimization for Flow Matching Models,” *arXiv preprint arXiv:2606.11025*, 2026.
- [10] Subliminal Clocks, “Latent Time Modelling in Diffusion Language Models,” *arXiv preprint arXiv:2607.01774*, 2026.